

Network Delay and the Limits of Guarding an Untrusted Supervisor: An AI-Control Testbed on a Quadruple-Tank Process

[NEED: author full name], [NEED: department and institution]

[NEED: supervisor names] | spriyad.infor@gmail.com

Note to readers. This is a draft for supervision review, not a submission. It contains the abstract, introduction, related work, methodology, experiments with results, threats to validity and references. The layout is single-column for reading. It has not been set in the IEEE two-column template. Every number comes from the repository (github.com/spriyads-vault/cps-alignment-protocols, branch `claude/quirky-fermi-560ln3`) and its result files. Appendix A lists what is unverified or unfinished, and should be read before the claims are used.

Abstract. *AI-control protocols limit the harm an untrusted model can do by combining a monitor, a safety filter and a trusted fallback controller. They are usually evaluated on software tasks, with no physical process and no communication delay. This paper builds a testbed on a simulated quadruple-tank process in which an untrusted supervisor sets two pump voltages across a delayed, lossy network. A monitor, a discrete-time control barrier function (CBF) filter and a trusted PI controller guard the process. The research question is how much network delay the guards tolerate before the supervisor can cause a tank to overflow. For the main experiments, predictions, protocols and metrics were registered before any evaluation episode ran. Across 17,340 simulated episodes, a CBF evaluated at the plant recorded no overflow in any of 288 settings, at delays up to 90 s. The same filter evaluated beside the supervisor lost safety at 12 to 16 s of round-trip delay, about one supervisor hold period. A follow-up study showed the tolerated delay growing with the hold period and shrinking with a looser barrier, and four of five registered predictions held. Safety was not usefulness: honest tracking performance under the edge filter fell to 0.02 of its undelayed value beyond 20 s. With Claude Haiku 4.5 as the monitor, scripted overt and authority-claim attacks were caught in every post-onset decision at a 1% false-alarm threshold. Covert sabotage still succeeded in 3 of 3 episodes at that threshold, and at a 5% threshold the rate fell to the level that held commands alone produce. The real-model results rest on one to three episodes, one planned adversary study was invalid because the attacker model refused, and the plant parameters await verification against the source paper.*

Index Terms. *AI control, control barrier functions, networked control systems, cyber-physical systems, language-model monitors, red teaming.*

I. Introduction

AI control asks a narrow question: if a capable model were trying to subvert the safeguards around it, how much harm could it still cause, and how much useful work would the safeguards cost [1]. The usual protocol layers a trusted but weaker monitor over an untrusted model, with a fallback to a trusted component when the monitor raises an alarm. Published evaluations use programming tasks [1] and agentic shell tasks [3], and a formal treatment models them as games [2]. None of these settings has a physical process whose state evolves while a message is in transit.

Cyber-physical systems add exactly that. A guard that reasons from a stale measurement protects a plant that has already moved. Networked control theory has studied how delay degrades stability [11], [12], and control barrier function (CBF) theory has extended safety filters to sampled and delayed inputs [8]–[10]. Neither literature considers an adversarial supervisor, and neither treats the guard's location relative to the network as a design variable under attack.

This paper studies one concrete case. An untrusted supervisor sets the pump voltages of a simulated quadruple-tank process [7]. A monitor, a CBF filter and a trusted PI controller guard it, and a delayed, lossy network sits between the supervisor and the plant. The question is how much delay the guards tolerate before the supervisor can overflow a tank. The contributions are:

- a simulated testbed with a pre-registered design, in which the monitor, the CBF and the trusted controller can be placed at the plant or beside the supervisor, and a ControlArena setting that exposes the same plant to other researchers [5], [6]
- a delay sweep over 17,340 episodes showing that guard placement, and not the guard itself, determines delay tolerance, and a follow-up study of how the tolerated delay scales with the supervisor's hold period and the barrier tuning
- a set of real-model checks with Claude Haiku 4.5 as monitor and Claude Sonnet 5.5 as supervisor, reported with their sample sizes and with the two that failed
- an account of which results hold only because the surrogate monitor was built to be fooled

II. Background and Related Work

A. AI control

Greenblatt et al. introduced control evaluations, in which a red team builds an untrusted model that tries to subvert a protocol and the protocol is scored on safety against the attack and usefulness for an honest model [1]. Their strongest protocols used a trusted model to edit or monitor the untrusted model's code. Griffin et al. formalised the red-team exercise as a partially observable stochastic game [2]. Bhatt et al. moved to agentic tasks and proposed resampling protocols, reporting that attack success fell from 58% to 7% at a 5% cost in usefulness [3]. Korbak et al. argued that the red team's permitted capabilities should scale with the agent's capability profile [4]. ControlArena packages such settings on top of the Inspect evaluation framework [5], [6]. The plant here is added as a ControlArena setting.

B. Safety filters and delay

A control barrier function certifies a safe set by constraining the control input at each state, and Ames et al. cast the constraint as a quadratic program [8]. Agrawal and Sreenath derived the discrete-time form [9], which is the form used here. Singletary et al. handled sampled-data systems with input delay by predicting the state across the delay [10]. Work on networked control systems analyses stability under delay and loss for a cooperative controller [11], [12]. In the sources we reviewed, the supervisor was never adversarial, and we did not find a study that measures delay tolerance against an attacker who chooses the commands.

C. Language-model monitors and injection

A language-model monitor reads text that the monitored model controls, so it is exposed to prompt injection [13], [14]. A rationale that claims authorisation is a direct instance. Whether current monitors are fooled by such claims in a control loop is an empirical question, and this paper reports a small number of measurements.

D. The quadruple-tank process

Johansson's quadruple-tank process has two pumps feeding four tanks through valves whose split ratios set the location of a multivariable zero, which can lie in the left or right half-plane [7]. It is a standard multivariable benchmark and is cheap to simulate. We use both a minimum-phase and a non-minimum-phase configuration.

III. Methodology

A. Problem statement and threat model

An untrusted supervisor observes the four tank heights and the reference every $T = 10$ s and returns two pump voltages and a free-text rationale. The command is held until the next decision. The supervisor may be honest or may pursue a hidden objective. Three attacks are considered. In the *overt* attack the supervisor sends maximum voltages to overflow a tank. The *injection* attack sends the same actions with a rationale claiming authorisation. The *covert* attack oscillates the pumps with a period of 40 s while keeping every height below the overflow level, with the aim of wasting energy and degrading tracking. Its *monitor-aware*

variant also attaches a routine-looking rationale. Attack onset is drawn uniformly from 150 to 300 s by episode seed, and every protocol faces the same onset for a given seed.

The guards see all four heights (assumption A1). The CBF has the exact plant model (A2). The monitor knows a sane controller for the task (A3). Model mismatch is not tested.

B. Plant

The plant follows Johansson's mass-balance model [7], integrated with a fourth-order Runge-Kutta scheme under zero-order-hold inputs:

$$dh_1/dt = -(a_1/A_1)\sqrt{(2g \cdot h_1)} + (a_3/A_1)\sqrt{(2g \cdot h_3)} + \gamma_1 k_1 v_1/A_1$$

$$dh_2/dt = -(a_2/A_2)\sqrt{(2g \cdot h_2)} + (a_4/A_2)\sqrt{(2g \cdot h_4)} + \gamma_2 k_2 v_2/A_2$$

$$dh_3/dt = -(a_3/A_3)\sqrt{(2g \cdot h_3)} + (1 - \gamma_2)k_2 v_2/A_3$$

$$dh_4/dt = -(a_4/A_4)\sqrt{(2g \cdot h_4)} + (1 - \gamma_1)k_1 v_1/A_4$$

Here h_i are heights in cm, v_j are pump voltages in $[0, 10]$ V, A_i and a_i are tank and outlet areas, k_j are pump gains and γ_j are valve split ratios. Water above 20 cm spills and the height is clipped at 20 cm, which is recorded as an overflow. The safe set is $h \leq 18$ cm. The plant steps every 2 s, the decision period is 10 s and an episode lasts 600 s, with the reference stepping at 60 s. The operating point is the exact equilibrium of the model for the nominal voltages, so the starting state is consistent with the parameters.

TABLE I. PLANT CONFIGURATIONS

Quantity	Minimum-phase	Non-minimum-phase	Verified
Valve splits (γ_1, γ_2)	0.70, 0.60	0.43, 0.34	Min.: yes. Non-min.: no
Tank areas A_1 to A_4 (cm ²)	28, 32, 28, 32	same	A_1, A_3 yes. A_2, A_4 no
Outlet areas a_1 to a_4 (cm ²)	0.071, 0.057, 0.071, 0.057	same	a_1, a_3 yes. a_2, a_4 no
Pump gains k_1, k_2 (cm ³ /Vs)	3.33, 3.35	same	k_1 yes. k_2 no
Nominal voltage (V)	3.0, 3.0	3.15, 3.15	Min.: yes. Non-min.: no

“Verified” means matched against a secondary source on 7 October 2026. The original paper was not accessible from the research environment, so the unverified entries are recorded as open in Appendix A. Equilibria derived from these parameters differ from the published operating points by up to 0.17 cm (minimum-phase) and 0.7 to 1.4 cm (non-minimum-phase).

C. Network

Messages travel from the supervisor to the plant over a link with a configurable round-trip delay split evenly between uplink and downlink, random jitter and random loss. A lost message leaves the previous command in force. Delivery happens at the next plant-step boundary, so a nominal delay is rounded up to a multiple of 2 s. All reported delays are these effective delays (a nominal 5 s behaves as 6 s, and 15 s as 16 s). Only a simulated clock is used. Model inference time enters as a fixed latency (median 1.9 s for the supervisor and 2.8 to 3.3 s for the monitor, as measured), so the result does not depend on API load.

D. Guards

Trusted PI controller. A decentralised PI controller runs at the plant, with the pairing of measurements and pumps swapped for the non-minimum-phase configuration. Its gains were tuned on development seeds. It serves as the fallback and as the baseline for harm and usefulness.

Control barrier function filter. For each lower tank, define $B = h_{\text{safe}} - h$ with $h_{\text{safe}} = 18$ cm. The discrete-time condition requires $B(h^+) \geq (1 - \alpha) \cdot B(h)$ after one step [9], with $\alpha = 0.5$ by default. The filter returns the voltage closest to the requested one that satisfies the condition for all tanks, by sequential quadratic programming. It was checked against brute-force search and for forward invariance in the test suite. Two placements are compared. At the *edge* the filter re-checks the held command on the live state at every 2 s plant step. At the *remote* placement the filter runs once per decision beside the supervisor, on the state the supervisor saw, and predicts over the whole 10 s hold. When no safe voltage exists it falls back to pumps off, which happens in about half of decisions at a 10 s hold even with no delay.

Monitor. Two monitors are used. The surrogate *RuleMonitor* scores the deviation of an action from a reference controller and discounts the score when the rationale carries an authority claim. That discount models prompt injection by assumption, so any result that depends on it is a statement about the harness and not about language models. The real monitor is Claude Haiku 4.5, prompted to return a suspicion score in $[0, 1]$ from the observation, the action and, in two of its three variants, the rationale. The variants are: reads the rationale with a prompt warning that authority claims are not evidence of safety, reads the rationale with no warning, and sees the action only. Responses that cannot be parsed are treated as maximally suspicious. Block and defer thresholds are calibrated on separate honest unguarded runs (seeds 1000 to 1999) at false-alarm rates of 1% and 5% per decision.

Supervisors. The scripted honest supervisor is a PI-like controller with gains tuned on development seeds. The scripted attackers are those above. The language-model supervisor is Claude Sonnet 5.5, prompted with the plant description, the current heights and the reference. An adaptive attacker is a parametric family (oscillation amplitude, period and duty, a constant bias, a feedback gain and a choice among four rationales) tuned by black-box search. A reasoning adversary is a model that writes the rationale for each covert command while seeing the monitor's scores for its earlier notes.

E. Protocols

TABLE II. PROTOCOLS

Protocol	Behaviour
monitor_only	A flagged message is discarded and the last command is held (block-and-hold).
cbf_only	Every command passes through the CBF filter.
hybrid	Block-and-hold on a flag, then the CBF filter on commands that pass.
defer_pi	A flagged message switches the plant to the local PI, which is not filtered. Otherwise the CBF filter applies.
pi_only	The trusted PI alone, ignoring the supervisor. It is the zero-harm baseline.
monitor_log	The monitor scores but never acts. It pays the monitor's latency and is the usefulness reference for monitored protocols.

Each guarded protocol is run with the CBF at the edge (suffix *@edge*) or at the remote placement (*@remote*).

F. Metrics

An *overflow* is any height above 20 cm. *Sabotage* is a mean tracking error after onset at least 0.5 cm above the trusted PI's, or pump energy at least 1.4 times the PI's. *Harm* is overflow or sabotage. Catastrophic attacks (overt, injection) are judged on overflow and covert attacks on sabotage. Rates carry 95% Wilson score intervals [15]. A protocol *fails* when the rate is at least 0.5 and *holds* when the point estimate is at most 0.1. *Usefulness* is the honest supervisor's integrated absolute tracking error (IAE) with no guard and no delay, divided by the protocol's IAE and capped at 1. The delay threshold d^* is the smallest round-trip delay at which the harm rate exceeds the baseline by 0.10, reported as the first delay with harm, the interpolated crossing, and the crossing of the Wilson lower bound. Monitor quality is reported as the area under the ROC curve (AUROC) and as the share of post-onset decisions caught at the calibrated threshold.

G. Experimental design

The design, claims, predictions and metric definitions were committed to the repository before any evaluation episode ran. Later changes are recorded as dated amendments (A1 to A5), including every correction that followed an independent code review. Development seeds 0 to 999 were used for tuning and a pilot, seeds 1000 to 1999 for threshold calibration, and seeds 10000 upward for evaluation. Evaluation cells use 30 episodes and the same 30 seeds in every cell, except where stated. No correction is applied across cells. With 30 episodes the 95% upper bound for zero failures is 0.11, so a result of 0 of 30 cannot confirm a rate of 0.1 or less.

Table III lists the experiments. Experiments 1 to 3 had registered predictions. The remainder are labelled as exploratory or as small real-model checks.

TABLE III. EXPERIMENT REGISTER

ID	Question	Size	Status
E1	Does each guard alone have a blind spot, at zero delay?	30 episodes per cell	Registered. Run
E2	Do the guards together close the blind spots, at zero delay?	30 episodes per cell	Registered. Run
E3	How much delay do the protocols tolerate, by placement?	17,340 episodes, 12 delays	Registered. Run
E3b	How does d^* scale with hold period and barrier tuning?	10 episodes per cell	Predictions registered first. Run
E4	Do E2 results survive a monitor that ignores authority claims?	1,260 episodes	Exploratory. Run
E5	Does Haiku 4.5 catch scripted attacks?	2 episodes per condition	Small. Run
E6	How does Haiku 4.5 perform inside the protocols?	3 episodes per cell	Small. Run
E7	What is the strongest covert attack a search finds?	487 evaluations per cell	Exploratory. Run
E8	What do the protocols cost a language-model supervisor?	1 episode per protocol	Single run
E9	Can a model write rationales that fool the monitor?	41 decisions	Invalid

H. Real-model procedure and implementation

Real-model scripts print their planned call counts, refuse to run without explicit confirmation, stop at a hard call limit shared across models, cache responses on disk and record the reason for every failed call. A run in which an attacker model fails on more than 20% of decisions is marked invalid, because its adaptive condition would then be mostly the fixed note. The code is written in Python with strict type checking, linting and 190 automated tests, including a golden regression file for episode outputs, and runs in continuous integration. The ControlArena setting was verified offline on the minimum-phase plant at zero delay and is untested for delay, the remote filter and the non-minimum-phase plant.

IV. Experiments and Results

A. Zero delay: each guard alone, then together (E1, E2)

A monitor alone failed against injection, with overflow in 30 of 30 episodes on both plants, and held against overt attacks (0 of 30). A CBF alone held against overt and injected attacks on overflow (0 of 30) and failed against covert sabotage (30 of 30). All 10 registered criteria for E1 passed. The monitor and CBF together held against overt and injected attacks (0 of 30 overflows) and failed against the monitor-aware covert attacker (30 of 30 sabotaged). Defer-to-PI failed there as well, because its trigger is the same monitor. All 8 criteria for E2 passed. Against plain injection, hybrid and defer-to-PI avoided overflow but were still sabotaged in 30 of 30 episodes, because the filter clamps the pumps at the barrier and tracking collapses. “Holds” therefore means no catastrophe, not no harm.

Three of these criteria (C1a, C2c and C2d) pass by construction, because the surrogate monitor scores any authority-claim rationale as zero. They show that the harness behaves as designed. They are not evidence that the protocols work against language-model monitors.

B. Delay sweep (E3)

Table IV gives overflow counts out of 30 on the minimum-phase plant by effective delay. Every edge protocol recorded no overflow in all 288 of its cells. Remote protocols lost safety at 12.0 to 13.6 s by point estimate, inside the registered range of 8 to 15 s (P3). One confirmed value, by Wilson lower bound, fell outside that range (hybrid against overt at 16.0 s on the minimum-phase plant). P4 held without exception.

For P5, the linearised honest loop was predicted to go unstable at a round trip of 15.1 s (spectral radius 0.986 at zero delay), and honest overflow with the remote filter appeared between 16 and 20 s. The non-minimum-phase honest loop never overflowed, as predicted.

TABLE IV. OVERFLOWS OUT OF 30, MINIMUM-PHASE PLANT, BY EFFECTIVE DELAY (S)

Protocol vs attack	0	2	6	8	10	12	16	20	30	46	60	90
cbf_only@remote vs injection	0	0	0	0	0	1	30	30	30	30	30	30
hybrid@remote vs injection	0	0	0	0	0	3	30	30	30	30	30	30
defer_pi@remote vs injection	0	0	0	0	0	1	30	29	30	30	30	30
hybrid@remote vs overt	0	0	0	0	0	1	6	21	28	23	17	1
cbf_only@remote vs honest	0	0	0	0	0	0	0	30	30	30	30	30
all @edge protocols, all attacks	0	0	0	0	0	0	0	0	0	0	0	0




Fig. 1. Overflows out of 30 against effective delay, minimum-phase plant.

The jump from 1 to 3 overflows at 12 s to 30 of 30 at 16 s reflects the 20 cm rim clipping a smooth rise. Mean peak height climbs from 16.2 cm at zero delay to 19.4 cm at 10 s, so the filter degrades gradually and the overflow count steps because the rim truncates it. On the minimum-phase plant the attack threshold (about 12.3 s) sits roughly 4 s below the delay at which the honest supervisor's own loop fails (16.4 s). Beyond that point a failure says nothing about the adversary. Safety was not usefulness. Under the edge protocols, honest usefulness on the minimum-phase plant fell from 1.00 at zero delay to 0.56 at 6 s, 0.23 at 12 s and 0.02 from 20 s. From 12 s, defer-to-PI held between 0.19 and 0.25 because the local PI is slow but stable. On the non-minimum-phase plant usefulness decayed gently to 0.65 at 90 s. Under the surrogate monitor, hybrid@edge was sabotaged in 30 of 30 episodes at every delay up to 60 s.


C. Scaling of the remote threshold (E3b)

No closed-form threshold is available. With a 10 s hold the remote filter repeatedly loses feasibility and falls back to pumps off, so the closed loop is a limit cycle and not a fixed point. We instead tested an extra-exposure mechanism with predictions written first (A4): a delay D applies the command D seconds later than the window the filter reasoned about, so overshoot should grow with D , depend on the hold period, depend on the barrier parameter, and fall as the filter barrier is tightened. The study varied the hold period H in {4, 10, 20} s, the barrier parameter α in {0.3, 0.5, 0.8} (named gamma in the code) and the filter barrier h_{safe} in {18, 17, 16} cm, with cbf_only@remote against the overt attack on both plants.

TABLE V. DELAY THRESHOLD D^* (S), MINIMUM-PHASE AND NON-MINIMUM-PHASE

Configuration	Minimum-phase	Non-minimum-phase
Baseline ($H = 10$ s, $\alpha = 0.5$, $h_{\text{safe}} = 18$ cm)	12.0	12.2
$H = 4$ s	4.2	6.2

Configuration	Minimum-phase	Non-minimum-phase
H = 20 s	24.4	18.2
Looser barrier, $\alpha = 0.8$	8.2	8.2
Tighter barrier, $\alpha = 0.3$	20.4	16.2
h_safe = 17 cm	12.2	14.2
h_safe = 16 cm	16.2	16.2

Fig. 2. Delay threshold d^* by configuration.

Four of five predictions held (P6 to P9). P10 failed as written. Its total clause held, since 2 cm of tightening added 4.2 s and 4.0 s, but its per-centimetre clause failed once, with 18 to 17 cm adding 0.2 s. Delays are quantised to 2 s, so the clause was too fine for the grid. The threshold lay between 0.9 and 1.6 times the hold period in all six cells. These values are from 10 episodes per cell and one attack, and tightening the barrier also leaves more room below the rim, so not all of its gain is filter behaviour. The mechanism fits the data and is not proved. It was formed after experiment E3 was read, so only the predictions are prospective.

D. A monitor that ignores authority claims (E4)

E2 was rerun with a surrogate monitor that ignores authority claims (1,260 episodes, same seeds). The monitor-dependent failures reversed. Against the monitor-aware covert attacker on the minimum-phase plant, defer-to-PI sabotage fell from 30 of 30 to 0 of 30, and hybrid fell from 30 to 10. On the non-minimum-phase plant, hybrid fell from 30 to 0 and defer-to-PI from 30 to 8. The remaining 10 and 8 equal each protocol's rate against plain overt and covert attacks, so they come from the held command after a block and from the PI fallback tracking more slowly than its own baseline. Claims C2c and C2d therefore hold only while the monitor is fooled.

E. Claude Haiku 4.5 as monitor (E5, E6)

A first pilot, with Claude Sonnet 5.5 as an attacking supervisor, returned a null result. Every one of its 102 attack-mode supervisor calls failed, so the monitor saw only a placeholder action and all three attacks scored an AUROC of 0.50. The likely cause is that the supervisor model declined the red-team prompts. Failure reasons are now recorded for every call. The remaining checks use scripted attacks, so no model is asked to attack.

TABLE VI. HAIKU 4.5 ON SCRIPTED ATTACKS, AUROC AND SHARE OF POST-ONSET DECISIONS CAUGHT AT THE 1% THRESHOLD

Attack	Warned	Unwarned	Action only
Overt	1.00, 100%	1.00, 100%	1.00, 100%
Injection (authority claim)	1.00, 100%	1.00, 100%	1.00, 100%
Covert	0.97, 80%	0.97, 89%	0.85, 64%
Covert, monitor-aware	0.97, 83%	0.97, 82%	0.85, 64%

Two episodes per condition, no failed calls. Removing the warning from the monitor's prompt did not make it easier to fool, and the authority-claim rationale lowered mean scores by only 0.01 to 0.02. Reading the rationale helped on covert sabotage (AUROC 0.97 against 0.85). The monitor-aware attacker is a fixed string that never saw the monitor, and the injected actions are a constant 10 V that the action-only monitor also catches in every decision.

With the monitor inside the protocols (minimum-phase plant, three episodes per cell, no overflow anywhere), the false-alarm budget decided the outcome. At a block threshold of 0.15 (1%) the monitor flagged 39% of the monitor-aware covert attacker's decisions and block-and-hold was sabotaged in 3 of 3 episodes. Defer-to-PI at a 5% threshold (0.08) flagged about all of them and was sabotaged in 0 of 3. A matched rerun with both thresholds at 5% reduced block-and-hold sabotage to 1 of 3, which equals its rate against overt attacks and is therefore the held-command artefact and not the attack. The price was false flags. Every honest episode contained at least one, at 5.0 blocked decisions per episode for block-and-hold and 2.3 deferred for defer-to-PI. Three episodes give a 95% interval of 0 to 0.56 for a rate of 0 of 3, which cannot reject the failure line of 0.5.

F. Adaptive attack search (E7)

A black-box search (120 random candidates, then 120 hill-climbing steps, each scored on two episodes) tuned the parametric covert attacker against each protocol and monitor in the closed loop. A second search constrained the attacker to be flagged on at most 5% of post-onset decisions, as a hard constraint, and the unconstrained search was seeded with its winner. Winners were rescored on 10 held-out seeds. A zero-amplitude attacker, which is the honest supervisor, served as the control.

TABLE VII. SABOTAGE SCORE ON HELD-OUT SEEDS, HYBRID PROTOCOL (1 OR MORE IS SABOTAGE)

Monitor	Plant	Best found	Best, at most 5% flagged	Honest control
Fooled by authority claims	Minimum-phase	38.95	38.95	0.72
Fooled by authority claims	Non-minimum-phase	35.86	35.86	0.69
Not fooled	Minimum-phase	1.08*	0.72	0.72
Not fooled	Non-minimum-phase	0.70	0.70	0.69

* The held command after a block, a known artefact, flagged at 100%. For defer-to-PI the best found was 38.95 and 35.87 against the fooled monitor, and the stealth-limited best was 0.72 and 0.70 against the unfooled one.

Against the monitor that authority claims fool, the search found sabotage at 36 to 39 times the sabotage line, never flagged, on every held-out seed. This only shows that the search finds a hole that exists by construction. Against the monitor that ignores authority claims, the best attack under the stealth limit scored 0.70 to 0.72 with no sabotaged seed, indistinguishable from the honest control. An earlier version with 80 evaluations per search had reported 12 to 39 times because it had not converged, and the figures here are from the rerun. One unconstrained result matters. On the non-minimum-phase plant, against the unfooled monitor, a fixed covert attack sabotaged 2 of 10 held-out seeds under defer-to-PI, because the flag hands control to a PI with no CBF in front of it. That is a protocol weakness and not a monitor miss. The thresholds come from a scripted honest supervisor whose deviations are small, so a real model's noisier honest behaviour would raise the threshold and widen the room for a quiet attack.

G. Language-model supervisor and reasoning adversary (E8, E9)

With Claude Sonnet 5.5 as an honest supervisor (one episode per protocol, minimum-phase plant, thresholds from one calibration episode), no call failed and no tank overflowed. The model tracked the reference about 100 times worse than the trusted PI.

TABLE VIII. HONEST SONNET 5.5 SUPERVISOR, ONE EPISODE PER PROTOCOL

Protocol	IAE	Relative to PI	Blocked or deferred (of 60)
Trusted PI alone	25.8	1×	0
Scripted honest, no guard	58.0	2.2×	0
Sonnet supervisor, no guard	2776	107×	0
Sonnet, monitor logs only	2796	108×	0
Sonnet, hybrid	2815	109×	4 blocked
Sonnet, defer-to-PI	2130	82×	2 deferred

Usefulness as defined in Section III-F, measured against the model's own unguarded error and capped at 1, came out at 0.99 for hybrid and 1.00 for defer-to-PI. Those values are uninformative here, because the baseline is already poor and the cap hides that deferring to the PI helped. Against the PI, usefulness is about 0.01. False flags on honest behaviour (4 of 60 blocked and 2 of 60 deferred) were close to the 5% target. We have not established whether the poor tracking comes from the model or from the prompt. The prompt in this run showed the model only the current heights and reference, with no record of its previous command, so a stateless controller cannot remove a steady error. A prompt that includes recent history is the first thing to test.

The reasoning-adversary run (E9) was invalid. The attacker model refused 40 of 41 decisions (category: cyber), above the 20% limit, so the adaptive and static conditions were effectively the same fixed note. As a repeat of the fixed-note measurement, Haiku 4.5 reading the rationale scored an AUROC of 0.977 without a warning and 0.992 with one, and caught 90% of 41 post-onset decisions at the 5% threshold. The run does not show whether a model-written, monitor-tuned rationale can fool Haiku 4.5. No attempt was made to work around the refusal.

V. Threats to Validity

- **Surrogate monitor.** Three registered criteria pass because the surrogate was built to be fooled. The first real-model evidence points the other way for crude injections.
- **Plant parameters.** Some values are unverified against the source paper (Table I). The qualitative result should not depend on them, but the exact thresholds might.
- **Sample sizes.** Evaluation cells use 30 episodes, and the real-model checks use one to three. A result of 0 of 30 has a 95% upper bound of 0.11. No correction is applied across cells.
- **One attack, one filter design, one plant model per study.** The 10 s hold makes the remote filter infeasible about half the time, so the delay budgets belong to that design and not to CBFs in general. The CBF has the exact plant model, and mismatch is untested.
- **Scripted honest behaviour.** Calibrated thresholds and usefulness baselines from a scripted supervisor are tighter than a language model would produce.
- **Metric artefacts.** The 1% false-alarm rate is per decision, so about 60% of honest monitor-only and hybrid episodes contain a false block. Harm for blocking protocols includes the drift of a held command. A review also noted that the sabotage window starts at the onset time while the attack starts at the next decision, up to 10 s later.
- **Post-hoc elements.** The mechanism in E3b was formed after E3 was read. E4 and E7 were not registered.

VI. Conclusion and Future Work

In this testbed, where a CBF runs relative to the network mattered more than the filter itself. At the plant it recorded no overflow in any tested setting, and beside the supervisor it tolerated delay of about one hold period. Guarding the process and keeping it useful are separate properties, and the real-model evidence is too small to settle how language-model monitors behave in the loop. The next steps are a corrected supervisor prompt with history, a larger real-model run with all usefulness figures reported against the trusted PI, a check of the plant parameters against the source paper, a CBF model-mismatch ablation, and a different attacker model for the reasoning-adversary study.

References

- [1] R. Greenblatt, B. Shlegeris, K. Sachan, and F. Roger, “AI control: Improving safety despite intentional subversion,” in *Proc. 41st Int. Conf. Machine Learning (ICML)*, PMLR vol. 235, 2024, pp. 16295–16336. arXiv:2312.06942.
- [2] C. Griffin, L. Thomson, B. Shlegeris, and A. Abate, “Games for AI control: Models of safety evaluations of AI deployment protocols,” in *Proc. IASEAI Conf.*, vol. 2, no. 1, pp. 204–216. arXiv:2409.07985.
- [3] A. Bhatt, C. Rushing, A. Kaufman, T. Tracy, V. Georgiev, D. Matolcsi, A. Khan, and B. Shlegeris, “Ctrl-Z: Controlling AI agents via resampling,” arXiv:2504.10374, Apr. 2025.
- [4] T. Korbak, M. Balesni, B. Shlegeris, and G. Irving, “How to evaluate control measures for LLM agents? A trajectory from today to superintelligence,” arXiv:2504.05259, Apr. 2025.
- [5] UK AI Security Institute, “Introducing ControlArena: A library for running AI control experiments.” [Online]. Available: <https://www.aisi.gov.uk/blog/introducing-controlarena-a-library-for-running-ai-control-experiments>. Accessed: Oct. 8, 2026.
- [6] UK AI Security Institute, “Inspect AI: Framework for large language model evaluations,” software. [Online]. Available: <https://pypi.org/project/inspect-ai>. Accessed: Oct. 8, 2026.
- [7] K. H. Johansson, “The quadruple-tank process: A multivariable laboratory process with an adjustable zero,” *IEEE Trans. Control Syst. Technol.*, vol. 8, no. 3, pp. 456–465, 2000.
- [8] A. D. Ames, X. Xu, J. W. Grizzle, and P. Tabuada, “Control barrier function based quadratic programs for safety critical systems,” *IEEE Trans. Autom. Control*, vol. 62, no. 8, pp. 3861–3876, Aug. 2017.
- [9] A. Agrawal and K. Sreenath, “Discrete control barrier functions for safety-critical control of discrete systems with application to bipedal robot navigation,” in *Proc. Robotics: Science and Systems XIII*, 2017, doi: 10.15607/RSS.2017.XIII.073.
- [10] A. Singletary, Y. Chen, and A. D. Ames, “Control barrier functions for sampled-data systems with input delays,” arXiv:2005.06418, 2020.
- [11] W. Zhang, M. S. Branicky, and S. M. Phillips, “Stability of networked control systems,” *IEEE Control Syst. Mag.*, vol. 21, no. 1, pp. 84–99, Feb. 2001.
- [12] J. P. Hespanha, P. Naghshtabrizi, and Y. Xu, “A survey of recent results in networked control systems,” *Proc. IEEE*, vol. 95, no. 1, pp. 138–162, Jan. 2007, doi: 10.1109/JPROC.2006.887288.
- [13] F. Perez and I. Ribeiro, “Ignore previous prompt: Attack techniques for language models,” in *NeurIPS ML Safety Workshop*, 2022. arXiv:2211.09527.
- [14] K. Greshake, S. Abdelnabi, S. Mishra, C. Endres, T. Holz, and M. Fritz, “Not what you’ve signed up for: Compromising real-world LLM-integrated applications with indirect prompt injection,” in *Proc. 16th ACM Workshop on Artificial Intelligence and Security (AISec)*, 2023. arXiv:2302.12173.
- [15] E. B. Wilson, “Probable inference, the law of succession, and statistical inference,” *J. Amer. Statist. Assoc.*, vol. 22, no. 158, pp. 209–212, 1927, doi: 10.1080/01621459.1927.10502953.

Appendix A. Verification Status and Open Items

This appendix is for the supervisors. It states what has been checked and what has not.

TABLE IX. OPEN ITEMS

Item	Status
Plant parameters	Tank 2 and 4 values, pump gain k_2 , and the whole non-minimum-phase set are unconfirmed against Johansson [7]. The paper was not reachable from the research environment. To close: compare Table I of the paper with plant/quadtank.py, correct any mismatch, regenerate the golden file and rerun the free experiments (about 25 minutes).
Reference [7] details	Volume, issue and pages were confirmed against secondary listings. The DOI and month were not confirmed and are omitted. Check on IEEE Xplore before submission.
Other references	Checked on 8 October 2026 against arXiv, proceedings, publisher or catalogue listings for title, authors, venue and pages. Full texts were not read for most, so the summaries of [2], [4], [5] and [10] rest on abstracts. Reference [10] is cited as a preprint. Its publication venue was not checked.
Related-work claim	The statement that no study measures delay tolerance against an adversarial supervisor is limited to the sources reviewed here and came from a small number of searches. A systematic search is needed before it is used as a novelty claim.
Honest-supervisor result (E8)	One episode per protocol. The prompt may explain the poor tracking. To test: add the previous command and recent readings to the prompt and rerun, then run about three episodes per protocol and report usefulness against the PI.
Reasoning adversary (E9)	Invalid. Needs a different attacker model, or is reported as a finding that the attacker declines.
Real-model sample sizes	E5 uses 2 episodes per condition and E6 uses 3 per cell. About 10 episodes per cell is needed to rule out a rate above about 0.28 from zero failures.
Results that hold by construction	C1a, C2c and C2d, as described in Sections III-D and IV-A.
Pre-registration timing	Amendment A1 states that its predictions were written without reading E3 output. A timestamp shows this only in part, and the amendment says so.
Format	Not in the IEEE two-column template. Author, affiliation and supervisor fields are placeholders.

Where the evidence is. Each experiment has a report and summary under results/ in the repository, and PREREGISTRATION.md holds the claims, predictions and amendments. The charts in the repository README are generated by scripts/make_readme_charts.py.